

# IEMS 5722

## Mobile Network Programming and Distributed Server Architecture

### Lecture 6

### Databases and Caches

Lecturer: Albert C. M. Au Yeung

18<sup>th</sup> February, 2016


# Data and Databases

Data can be considered as the **most important assets** in many Internet-based services, consider:

- The social network and users' interests in Facebook
- The tweets in Twitter
- The search index and cache in Google
- ...



# Three-Tier Architecture



# Data and Databases

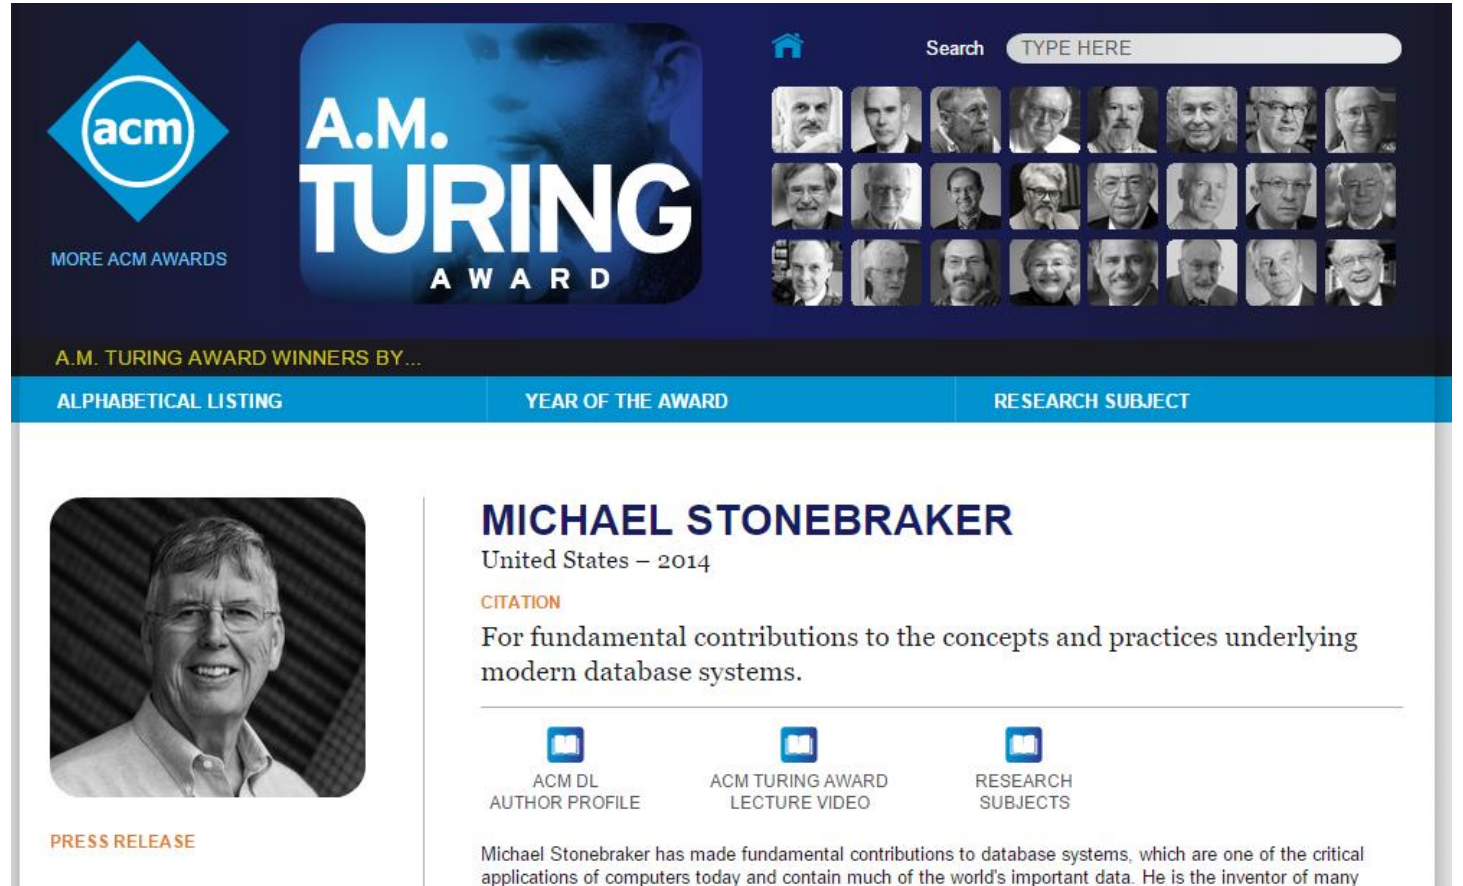
Most Internet-based services can be considered as some means for interacting with some data



# Turing Award 2014

## Michael Stonebraker

Involved in the invention and development of many relational database concepts (e.g. the object-relational model, query modification, etc.)



The screenshot shows the ACM Turing Award website. At the top left is the ACM logo with the text "MORE ACM AWARDS". To its right is a large graphic with "A.M. TURING AWARD" and a portrait of Alan Turing. On the top right is a search bar with a home icon and the text "Search TYPE HERE". Below the search bar is a grid of 24 small portraits of award winners. A navigation bar below the grid has three tabs: "ALPHABETICAL LISTING", "YEAR OF THE AWARD", and "RESEARCH SUBJECT". The main content area features a large portrait of Michael Stonebraker on the left, with the text "PRESS RELEASE" below it. To the right of the portrait, the name "MICHAEL STONEBRAKER" is displayed in large blue letters, followed by "United States – 2014". Below this is a "CITATION" section with the text: "For fundamental contributions to the concepts and practices underlying modern database systems." At the bottom of the profile section are three icons with labels: "ACM DL AUTHOR PROFILE", "ACM TURING AWARD LECTURE VIDEO", and "RESEARCH SUBJECTS". A paragraph at the very bottom states: "Michael Stonebraker has made fundamental contributions to database systems, which are one of the critical applications of computers today and contain much of the world's important data. He is the inventor of many".

[http://amturing.acm.org/award\\_winners/stonebraker\\_1172121.cfm](http://amturing.acm.org/award_winners/stonebraker_1172121.cfm)

# Relational Databases

# Database Management Systems

## Database Management System (DBMS)

- A system that stores and manages a (probably large) collection of data
- It allows users to perform operations and manage the data collection (e.g. creating a new record, querying existing records)
- Examples
  - › Oracle
  - › MS SQL Server
  - › MySQL
  - › Postgre SQL

# Database Management Systems

## Data Model

- A data model describes how data should be organised
- It describes how data elements relate to one another
- In most cases, a data model reflects how things are related in the real world

A widely used data model is the *relational model of data*

- A table describes a relation between different objects



# Relational Databases

- A database is a collection of relations (tables)
- Each relation has a list of attributes (columns)
- Each attribute has a domain (data type)
- Each relation contains a set of tuples (rows)
- Each tuple has a value for each attribute of the relation (or NULL if no value is given)

# Relational Databases

## Simple Example – Student Enrollment in Courses

Students

ID	Name	Year
1	John Chan	3
2	May Lee	4
...	...	...

Courses

ID	Code	Lecturer
1	IEMS 5723	...
2	IEMS 5722	...
...	...	...

Students Enrollment

ID	Student ID	Course ID
1	1	1
2	2	1
...	...	...

# Relational Databases – Schema & Instance

## Schema (also known as metadata)

- Specifies how data is to be structured
- Needs to be defined before the database can be populated

## Instance

- The actual content to be stored in the database
- The structure of the data must conform to the schema defined beforehand

# Relational Databases – Schema & Instance

## Example

- A table "Student" with the following schema
  - › (ID integer, name string, year integer, date\_of\_birth date)
- Some instances of "Student" in the table:
  - › (1, 'Peter Chan', 3, 1996-03-17)
  - › (2, 'Mike Cheung', 3, 1996-05-19)
  - › ...

# Relational Databases

How can we create schema and modify the data in a database management system?

## SQL – Structured Query Language

- A standard language for querying and manipulating data in a relational database
- It is both a DDL (data definitional language) and a DML (data manipulation language)
- Defining schemas with "**create**", "**alter**", "**delete**"
- Manipulating tables with "**insert**", "**update**", "**delete**"

# SQL Introduction

Let's assume we have the following two tables

Students

id	name	year
1	John Chan	3
2	May Lee	4
...	...	...

Courses

id	code	lecturer
1	IEMS 5723	...
2	IEMS 5722	...
...	...	...

Enrollment

id	student_id	course_id
1	1	1
2	2	1
...	...	...

# SQL Introduction

How can we create these tables?

```
CREATE TABLE Student (  
    id INT NOT NULL AUTO_INCREMENT,  
    name VARCHAR(100) NOT NULL,  
    year INT NOT NULL,  
    PRIMARY KEY (id)  
);
```

**AUTO\_INCREMENT:**  
Wherever you insert a new row into the table, it automatically increments by 1

**PRIMARY KEY:**  
A key of the table, it can be used to uniquely identify a particular record in the table

```
CREATE TABLE Courses(  
    id INT NOT NULL AUTO_INCREMENT,  
    code VARCHAR(10) NOT NULL,  
    lecturer VARCHAR(100) NOT NULL,  
    PRIMARY KEY (id),  
    UNIQUE (code)  
);
```

**UNIQUE:**  
The field must be unique for each row in the table

# SQL Introduction

## **SELECT statement**

- Used to retrieve data from one or more tables given some conditions
- Example 1: retrieve the E-mail address of the student 'John Chan'

```
SELECT email FROM Students WHERE name = 'John Chan';
```

- Example 2: retrieve the name of the lecturer of course 'IEMS 5722'

```
SELECT lecturer FROM Courses WHERE code = 'IEMS 5722';
```



# SQL Introduction

- Example 3: Retrieve a list of students whose name is 'John'

```
SELECT * FROM Students WHERE name LIKE 'John %';
```

- Example 4: Retrieve a list of courses, sort by their course code in descending order

```
SELECT id, code, lecturer  
FROM Courses  
ORDER BY code DESC
```

# SQL Introduction

- Example 5: Retrieve a list of students who have enrolled in 'IEMS 5722'

```
SELECT s.id, s.name
FROM Students s, Courses c, Enrollment e
WHERE
    e.student_id = s.id
AND e.course_id = c.id
AND c.code = 'IEMS 5722'
```



Here, we are **joining (inner join)** the three tables in order to retrieve data based on their relationships

## References:

- [https://en.wikipedia.org/wiki/Join\\_\(SQL\)](https://en.wikipedia.org/wiki/Join_(SQL))
- <http://blog.codinghorror.com/a-visual-explanation-of-sql-joins/>

# SQL Introduction

## **INSERT statement**

- Used to insert new data into the tables
- Example 1: Insert a new student into the Students table

```
INSERT INTO Students (name, email)  
VALUES ('Paul Wong', 'pw@gmail.com')
```

- Example 2: Insert a new course into the Courses table

```
INSERT INTO Courses (code, lecturer)  
VALUES ('IEMS 5678', 'Prof. Cheung')
```

# SQL Introduction

## **UPDATE statement**

- Used to modify the data in the tables
- Example 1: Change the email address of the student with id = 12

```
UPDATE Students  
SET email = 'abc123@gmail.com'  
WHERE id = 12
```

- Example 2: Update the lecturer of the course with course code 'IEMS 3456'

```
UPDATE Courses  
SET lecturer = 'Prof. Chan'  
WHERE code = 'IEMS 3456'
```

# SQL Introduction

## **DELETE statement**

- Used to modify the data in the tables
- Example 1: Change the email address of the student with id = 12

```
UPDATE Students  
SET email = 'abc123@gmail.com'  
WHERE id = 12
```

- Example 2: Update the lecturer of the course with course code 'IEMS 3456'

```
UPDATE Courses  
SET lecturer = 'Prof. Chan'  
WHERE code = 'IEMS 3456'
```

# SQL Introduction

For more complex SQL statements and queries, refer to the tutorials in the following Web sites

- MySQL Reference Manual:  
<http://dev.mysql.com/doc/refman/5.7/en/tutorial.html>
- MySQL Tutorial: <http://www.mysqltutorial.org/>
- W3School SQL Tutorial: <http://www.w3schools.com/sql/>

# ACID Properties of Relational Database

Relational databases focus on having reliable transactions, and usually have the ACID properties

- *Atomicity* – Each transaction is either “all done” or “failed”
- *Consistency* – Data can only be changed according to pre-defined rules
- *Isolation* – Concurrent queries do not interfere with one another
- *Durability* – Results are persistent in the databases

# MySQL



- An open source relational database management system
- The world's second most widely used RDBMS
- Most widely used RDBMS in a client-server model
- <http://www.mysql.com/>
- Community Edition – freely available on Windows, Mac OS and Linux
- Enterprise Edition – More advanced functions with technical support
- In Ubuntu, install the MySQL server with

```
$ sudo apt-get install mysql-server
```



# MySQL

- Once installed, you can use its command line client to interact MySQL

```
$ mysql -uroot -p
...
mysql> show databases;
+-----+
| Database |
+-----+
| information_schema |
| mysql |
| performance_schema |
| phpmyadmin |
+-----+
5 rows in set (0.01 sec)
```

```
mysql> create database iems5722;
Query OK, 1 row affected (0.03 sec)
```

```
mysql> use iems5722;
Database changed
```

```
mysql> show tables;
Empty set (0.00 sec)
```

# Interfacing MySQL in Python

# MySQL & Python

- In your server application, it is very likely that you will have to access or modify the data stored in the database
- In Python, you can make use of the MySQLdb module to help you execute SQL statements (<http://mysql-python.sourceforge.net/MySQLdb.html>)
- Install the MySQLdb module using the following command:

```
$ sudo apt-get install python-mysqldb
```

- Check if it has been installed successfully by importing MySQLdb

```
>>> import MySQLdb  
>>>
```

## Connecting to the MySQL database

```
import MySQLdb

db = MySQLdb.connect(
    host = "localhost",
    port = 3306,
    user = "dbuser",
    passwd = "password",
    db = "mydb",
    use_unicode = True,
    charset = "utf8",
    cursorclass = MySQLdb.cursors.DictCursor
)
```

## Executing an SQL query

```
query = "SELECT * FROM Students ORDER BY id ASC"
```

```
# Execute the query
```

```
db.cursor.execute(query)
```

```
# Retrieve all the results
```

```
results = db.cursor.fetchall()
```

```
# results is a list of rows, each is a dictionary
```

```
# The following line prints 'John Chan'
```

```
print results[0]['name']
```

# MySQL & Python

- You can also fetch records one after another

```
query = "SELECT * FROM Students ORDER BY id ASC"
```

```
# Execute the query
```

```
db.cursor.execute(query)
```

```
# Retrieve rows one by one
```

```
row = db.cursor.fetchone()
```

```
while row is not None:
```

```
    print row['name']
```

```
...
```

## Parameter substitution

- Very often you have values stored in Python variables, and would like to use them in the SQL queries

```
student_id = request.form.get("student_id")
email = request.form.get("email")

query = "UPDATE Students SET email = %s WHERE id = %s"

# prepare the parameters (must be a tuple!)
params = (email, student_id)

# Execute the query by substituting the parameters
db.cursor.execute(query, params)
db.commit() # Remember to commit if you have changed the data!
```

## Executing multiple queries

- Sometimes you may want to execute many queries with a list of values

```
students = [  
    ('May Chan', 'mc@gmail.com'),  
    ('Peter Lo', 'pl@gmail.com'),  
    ('William Wong', 'ww@gmail.com')  
]  
  
query = "INSERT INTO Students (name, email) VALUES (%s, %s)"  
  
# Execute multiple queries at a time with a list of parameters  
db.cursor.executemany(query, students)  
db.commit()
```



# Using MySQL with Your Flask App

# Connecting to MySQL in Flask

- Recall that we use Flask to develop our APIs for our mobile apps

```
from flask import Flask
app = Flask(__name__)

@app.route('/get_students')
def get_students():
    ...

if __name__ == '__main__':
    app.run()
```

# Connecting to MySQL in Flask

- What if we need to develop an API for retrieving the list of students from the database?

```
from flask import Flask
app = Flask(__name__)

@app.route('/get_students')
def get_students():
    # 1. Connect to database
    # 2. Construct a query
    # 3. Execute the query
    # 4. Retrieve data
    # 5. Format and return the data

if __name__ == '__main__':
    app.run()
```

# Connecting to MySQL in Flask

- For readability and reusability, let's create a class that will help us connect to the database

```
class MyDatabase:
    db = None

    def __init__(self):
        self.connect()
        return

    def connect(self):
        self.db = MySQLdb.connect(
            host = "localhost",
            port = 3306,
            user = "...",
            passwd = "...",
            db = "...",
            use_unicode = True,
            charset = "utf8",
            cursorclass = MySQLdb.cursors.DictCursor
        )
        return
```

# Connecting to MySQL in Flask

- Let's implement the get\_students() function

```
@app.route('/get_students')
def get_students():

    # Create the database object
    mydb = MyDatabase()

    # Prepare and execute the query
    query = "SELECT * FROM Students"
    mydb.db.cursor.execute(query)

    # Retrieve the data and send response
    students = mydb.db.cursor.fetchall()
    return jsonify(data=students)
```

# Connecting to MySQL in Flask

- Let's see another example, what if we need to implement an API for retrieving the data of a single student?

```
@app.route('/student/<int:student_id>')
def get_single_student():

    mydb = MyDatabase()
    query = "SELECT * FROM Students WHERE id = %s"
    params = (student_id,) # Note the comma here!
    mydb.db.cursor.execute(query, params)

    student = mydb.db.cursor.fetchone()
    if student is None: # No such student is found!
        return jsonify(status="ERROR", message="Not Found!")
    else:
        return jsonify(status="OK", data=students)
```

To use this API, send a GET request to, for example, **/student/2**

(Retrieve data of the student with id = 2)

## Before and After Request

- In Flask, you can specify some codes to be executed before and/or after a request from the client
- This is done by implementing the `before_request` and `teardown_request` functions

```
@app.before_request  
def before_request():  
    # Your code here...  
    return  
  
@app.teardown_request  
def teardown_request(exception):  
    # Your code here...  
    return
```

## Before and After Request

### **How would you use these two functions?**

1. Create a database connection before a request, and close the connection after the request
2. Log the request to the database or to a file before each request
3. Check user authentication before each request
4. ...



# Before and After Request

## Example

- We create the database connection before the request, store it in the globally available object 'g', and close the connection after the request

```
@app.before_request
def before_request():
    g.mydb = MyDatabase()
    return

@app.teardown_request
def teardown_request(exception):
    mydb = getattr(g, 'mydb', None)
    if mydb is not None:
        mydb.db.close()
    return
```

**g** is an object that is available throughout the whole request, thus it will be available to you in the API functions you implement

Remember to import it by:  
**from flask import g**

## Before and After Request

Then, in our API function, we can simply write:

```
@app.route('/student/<int:student_id>')
def get_single_student():

    query = "SELECT * FROM Students WHERE id = %s"
    params = (student_id,) # Note the comma here!
    g.mydb.db.cursor.execute(query, params)

    student = g.mydb.db.cursor.fetchone()
    if student is None: # No such student is found!
        return jsonify(status="ERROR", message="Not Found!")
    else:
        return jsonify(status="OK", data=students)
```

Reference: <http://flask.pocoo.org/docs/0.10/tutorial/dbcon/>

# NoSQL Databases

# NoSQL

The relational model of data and relational databases are powerful tools for managing data, but they cannot solve all problems

- Data Model - data may be better modelled as **objects in a hierarchy** or a **graph**
- Scheme - in many applications, it can be too restrictive to have **fixed schema**
- Scalability - it takes a lot of effort to **horizontally scale relational databases**


Alternative solutions are therefore desirable for solving new problems

# NoSQL

NoSQL (non-SQL, non-relational, not-only-SQL) systems are storage systems that offer users the ability to model data in ways other than relational tables.

- It is NOT a single technology
- No single definition of a NoSQL database
- Many different systems or solutions are available for solving different problems

## *HOW TO WRITE A CV*



# Why do we need NoSQL Databases?

## 1. Popularity of Web applications and services

- Many writes and reads because of user participation (user-generated content)
- Complex functions require flexibility in data models (e.g. find friends of friends in a social network, find related items bought by users of the same age group, ...)
- Horizontal scaling is desirable

# Why do we need NoSQL Databases?

## 2. Flexibility in data schema is required

- Relational database requires data schema to be well-defined
- However, in many applications there can be a lot of attributes and these attributes may change over time

## 3. Different solutions required to handle different types of data

- Structured vs. semi/unstructured data
- Data that needs to be served real-time vs. log data

# NoSQL

Some common features of NoSQL database systems:

- Do not require the definition of a fixed schema
- Scale horizontally (distributed operations, replication and partition) over multiple servers
- Simple or no query language, offer APIs for manipulating the data
- A weaker concurrency model (not ACID)
- Distributed storage



# NoSQL Database Systems


## The different types of NoSQL database systems

- Key-value stores
- Document databases
- Graph databases
- Column databases
- Object databases

# NoSQL Database Systems

## Key-value stores

- Examples are Redis, Riak, Oracle NoSQL Database
- Implementing a dictionary or a hash
- Retrieval of data is very fast
- For quickly retrieving the value of a known key, but not good for searching



# NoSQL Database Systems

## Document Stores

- Examples are CouchDB and MongoDB
- Similar to key-value stores, but value is a document
- Document is in a semi-structured format (e.g. JSON or XML)
- Allow retrieval of documents by searching their content

# NoSQL Database Systems

## Graph Databases

- Examples are Neo4j, Titan and OrientDB
- Store data in the form of
  - › Nodes (entities)
  - › Edges (relations between entities)
  - › Properties (attributes of nodes or edges)
- Perform queries on graphs without the need to carry out expensive JOIN operations

# NoSQL Data Models

Data models in NoSQL databases are very different from that in relational databases

Major principles in NoSQL data modelling are:

- Denormalisation
- Aggregation
- Application-level Join

# NoSQL Data Models

## Normalisation

- Database normalisation is the process of organising tables in a relational database to minimize data redundancy

Image ID	Image Name	User	Tag
1	...	A	Cat
2	...	A	Dog
3	...	B	Cat
4	...	B	Fish




Image ID	User Id
1	1
2	1
3	2
4	2

Image ID	Tag ID
1	1
2	2
3	1
4	3

Image ID	Image Name
1	...
2	...
3	...
4	...

# NoSQL Data Models

## Denormalisation

- Normalisation ensures minimal redundancy, but then you will need to perform (a lot of) join operations to get what you want
- Denormalisation is the opposite, to improve *performance* and *scalability*, we add redundant data such that we can avoid joins
- Very fast read, but may have more complex and slower write/update logic  
→ *not a problem because writes can wait*

# NoSQL Data Models

## Aggregation


- Because NoSQL databases usually do not place constraints on values and schemas, we can aggregate different objects and attributes under the same record for the sake of performance
- Especially when dealing with many-to-many relations (e.g. hashtags and images, tracks and albums, users' comments on one another, etc.)



# NoSQL Data Models

## Application-Level Join

- Combines data in the application instead of relying on complex join queries in the database



# NoSQL Data Models

## Application-Level Join

### Pros:

- No complex join operations on the database side, retrieval of data can be very fast
- Frequently used data can be cached for even faster retrieval

### Cons:

- Most likely you will have to issue multiple queries to the database to retrieve data
- Application may need to be optimised to ensure performance

# Redis

- <http://redis.io/topics/introduction>
- An open source in-memory data structure store
- Can be used as a key-value database, cache, or message broker
- Install redis in Ubuntu with the following command

```
$ sudo apt-get install redis-server
```

- You can check if the server has been installed successfully by running the redis command line tool:

```
$ redis-cli  
127.0.0.1:6379>
```

# Redis

- You can easily interface with Redis in Python
- Install the Python redis module with the following command:

```
$ sudo pip install redis
```

- Check whether the installation is successful:

```
>>> import redis  
>>>
```

# Redis

## A Simple Example

```
# import redis
from redis import StrictRedis

# Establish a connection to redis on localhost
r = StrictRedis('localhost')

# Set the value of a key
r.set('test_key', 'test_value')

# Get the value of a key
# value will be None if no such key is found in redis
value = r.get('test_key')
```

# Redis

- You can store strings, lists, sets, or even bit arrays in Redis
- It also supports counters (increment or decrement the value)
- More examples below:

```
# Create a counter, initialise it  
r.set('counter', 1)  
  
# Increment the counter  
r.incr('counter')  
...  
  
# Push a string into a list  
r.rpush('user_list', 'John Chan')  
...
```

References: <http://redis.io/topics/data-types-intro>

# Caching

# Caching

**Cache** is a temporary data storage that stores data for quick retrieval in the future

- Mostly implemented as a key-value store, where the unique key can be used to retrieve the value at  $O(1)$  time
- Cache is usually small (RAM is expensive!)
- Hit (found) vs. Miss (not found)
- Cache can be persistent, if it also stores the current state into some persistent storage (e.g. the hard disk)



# Caching

Where should you use cache?



Client

API Requests

Load  
Balancer


Application  
Servers

File Storage  
Servers

(Master)

(Slaves)

Database  
Servers



# Memcached

A general purpose distributed memory caching system

- *General purpose* – can be used in front of a Web server, an application server, or a database server
- *Distributed* – can be operated on multiple servers for scalability
- *Memory* – stores values in RAM, if not enough RAM, discard old values

# Memcached + Nginx

```
server {  
    location / {  
        set $memcached_key "$uri?$args";  
        memcached_pass host:11211;  
        error_page 404 502 504 = @fallback;  
    }  
  
    location @fallback {  
        proxy_pass http://backend;  
    }  
}
```

The key-value pair should be inserted into Memcached by the application (external to Nginx)

# Memcached + MySQL

```
import sys
import MySQL
import memcache

mem = memcache.Client(['127.0.0.1:11211'])
conn = MySQLdb.connect(...)
...

user_record = memc.get('user_5')

if not user_record:
    # retrieve user record from MySQL
else:
    # data available, retrieved from Memcached
```

# Memcached

More references can be found at:

- <http://memcached.org/>
- Python Memcached module:  
<https://pypi.python.org/pypi/python-memcached>
- Caching in Flask:  
<http://flask.pocoo.org/docs/0.10/patterns/caching/>
- Using MySQL and Memcached with Python:  
<https://dev.mysql.com/doc/mysql-ha-scalability/en/ha-memcached-interfaces-python.html>

# Next Lecture:

## Instant Messaging and Push Notifications

(Create a Google account if you have not)

End of Lecture 6