

# IEMS 5722

## Mobile Network Programming and Distributed Server Architecture

### Lecture 9

### Asynchronous Tasks & Message Queues

Lecturer: Albert C. M. Au Yeung

10<sup>th</sup> March, 2016

# The HTTP Request-Response Cycle

- The application server only performs work whenever there is a *request* from the client
- The client *waits* for the response before the application server has completed its work and returns a *response*



# The HTTP Request-Response Cycle

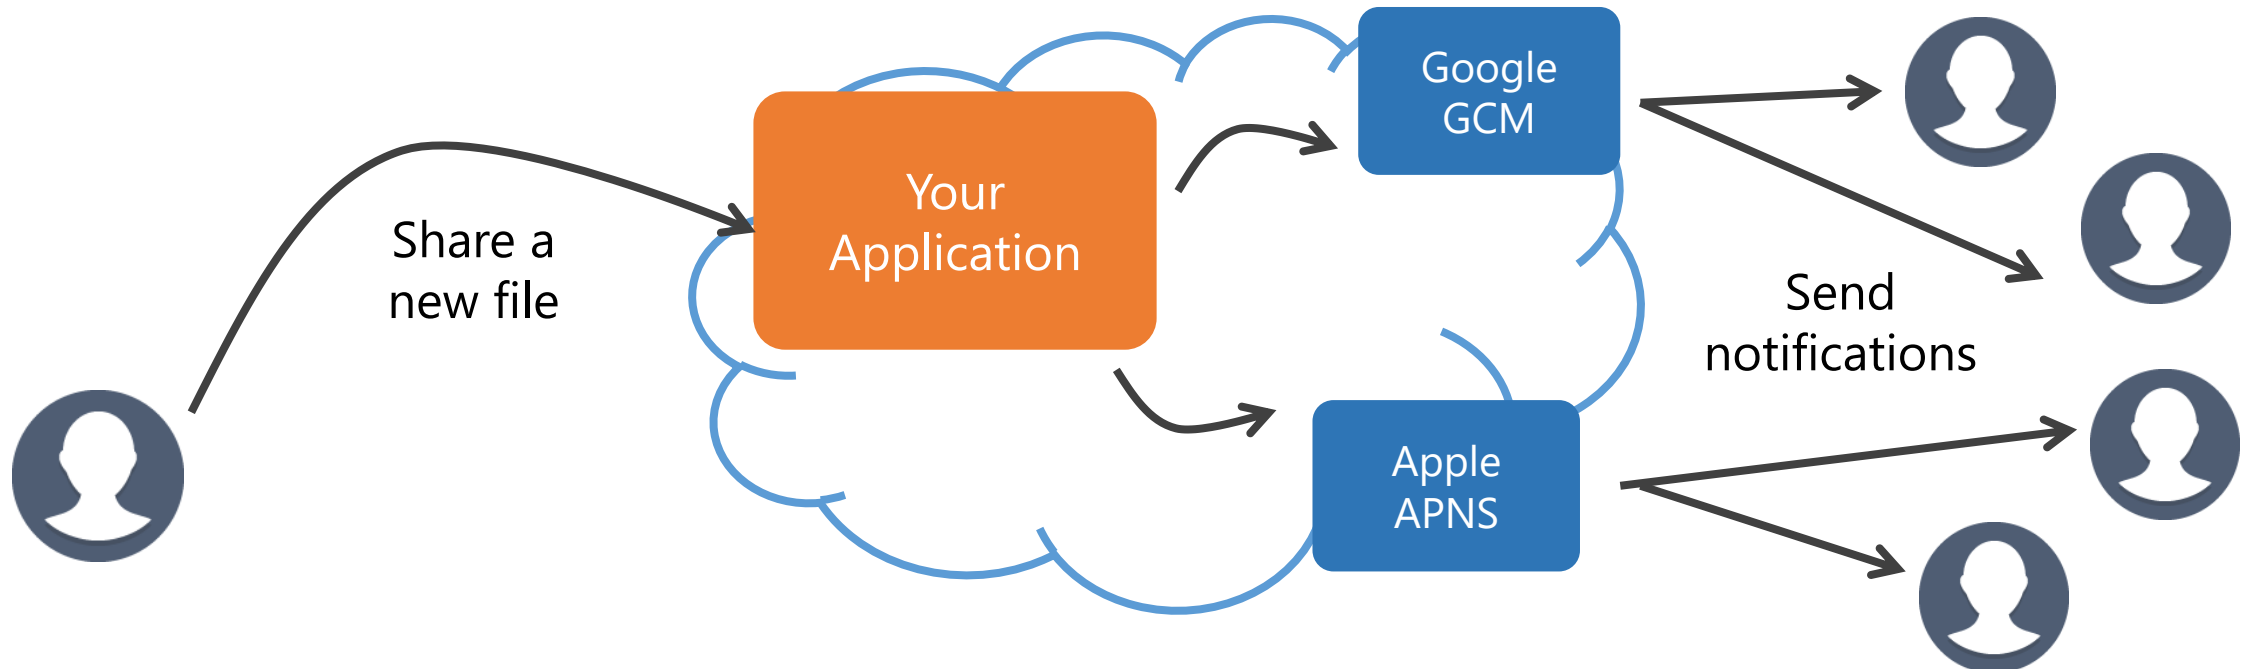
- The **HTTP request-response cycle** is expected to complete in a *short* time (no one likes waiting!)
- However, not all tasks can be completed in a short time
- Therefore, it is necessary to carry out some tasks in the background (i.e. *outside* the HTTP request-response cycle)



## Example

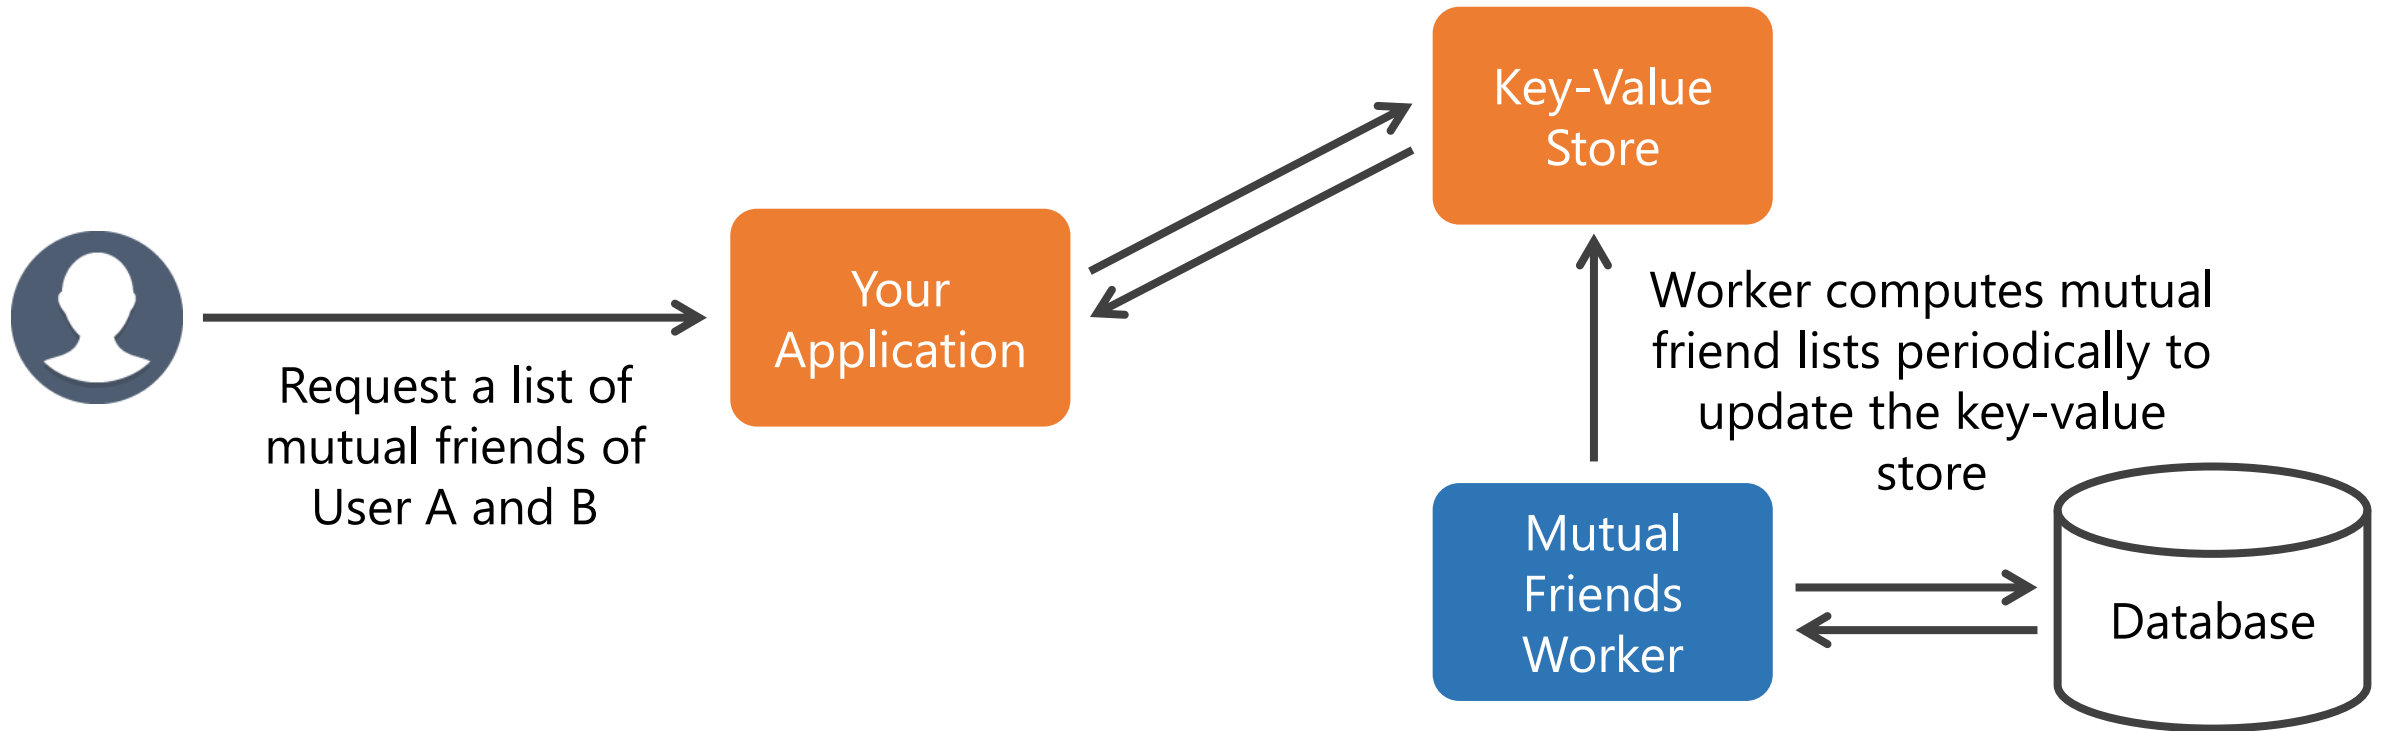
There are a lot of scenarios in which background tasks are necessary

- Consider a mobile app in which **notifications** will be sent to your friends when you have shared a new file



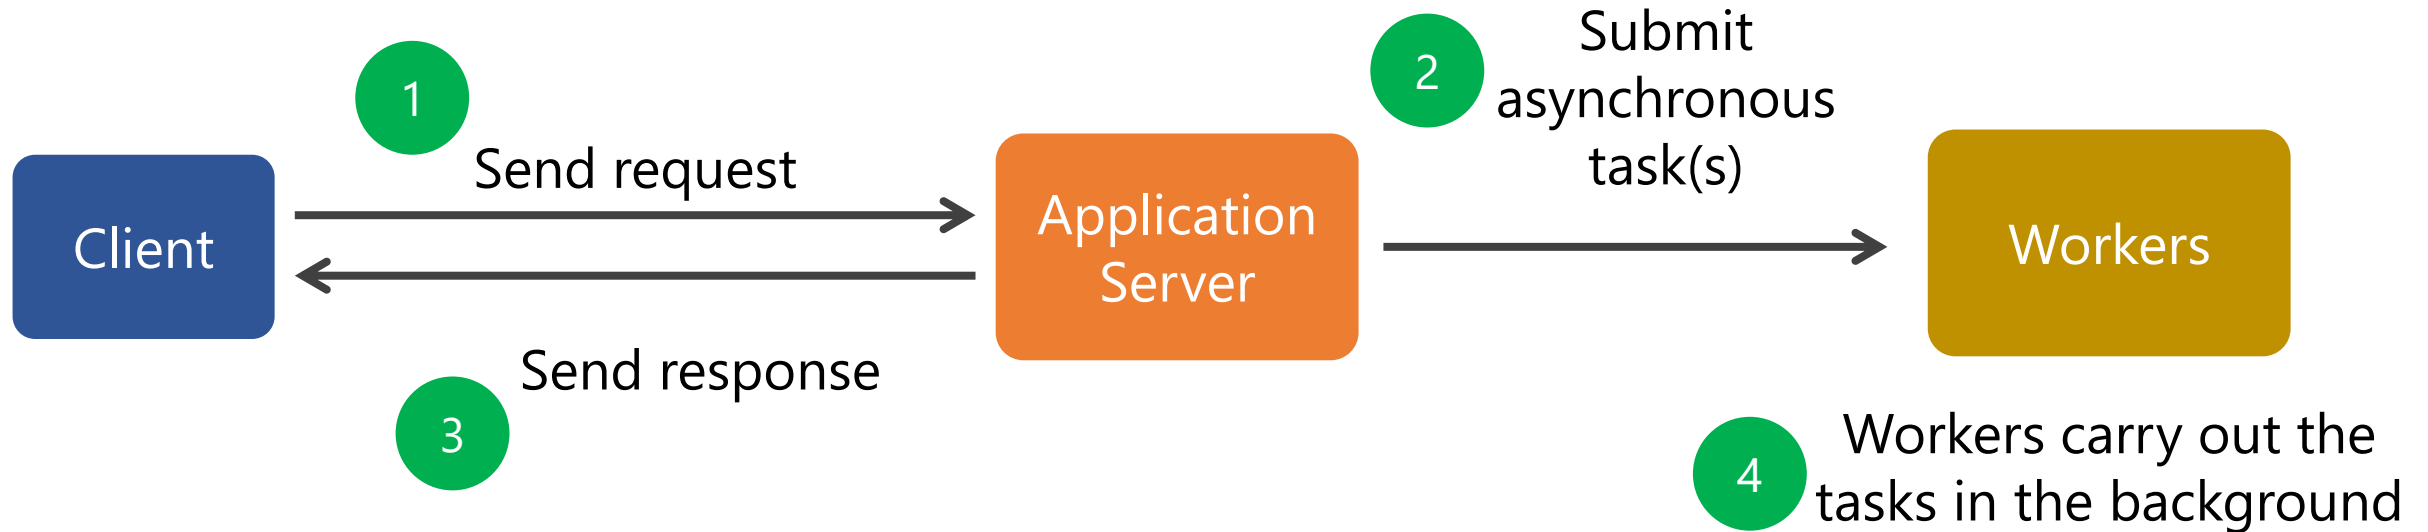
## Example

- Consider a social network in which you need to quickly retrieve the list of mutual friends of two users



# Asynchronous Tasks

- Task carried out outside the HTTP response-request cycle are called *asynchronous (non-blocking) tasks*
- Achieve de-coupling by separating the application server and other services



## Regular Jobs

A simple way to implement asynchronous tasks is to schedule “jobs” in the system to be run at specific time (and periodically)

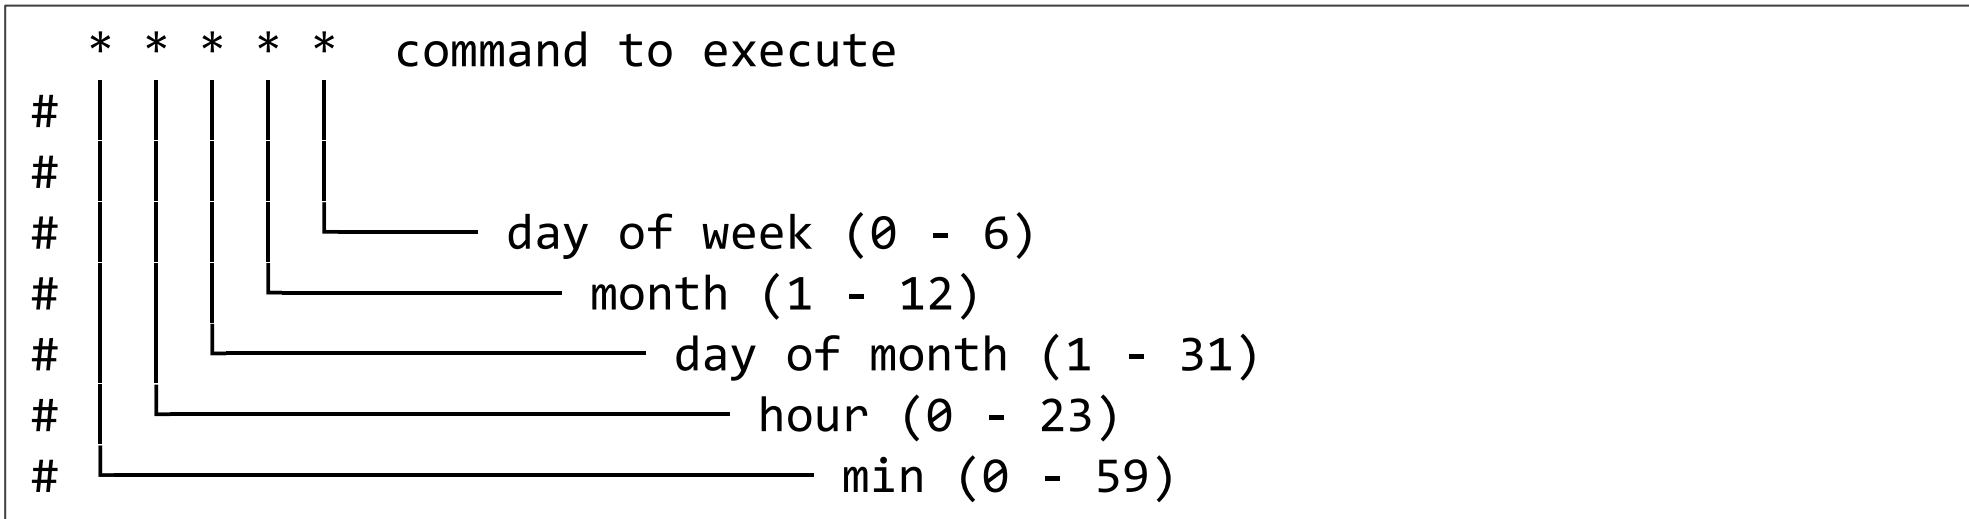
For example:

- Update various counters in the system (e.g. number of users, posts, messages)
- Generate a list of recommended friends for a user
- Compute the similarity of two products based on user feedbacks in a e-commerce site
- Collect information from an external source (e.g. news or weather info)

# Cron

**Cron** is a scheduler in Unix/Linux for setting up regular jobs

- Type “**crontab -e**” to edit the configuration file (**cron table**)
- Each line in the configuration file defines a single job





# Cron Examples

- Run at 10:30am everyday

```
30 10 * * * /home/user/program.py
```

- Run at the 0<sup>th</sup> minute of every hour

```
0 * * * * /home/user/program.py
```

- Run at 6pm on every Wednesday

```
0 6 * * 3 /home/user/program.py
```

- Run every 10 minutes on Monday to Friday

```
*/10 * * * 1-5 /home/user/program.py
```

## Limitation of Cron Jobs

Cron is used to schedule periodic tasks. However:

- Requests sent to your server may not be evenly distributed across time
- Not all tasks can be finished within a certain amount of time
- It is not trivial to incorporate logic or load balancing in carrying out a set of related jobs

And there are other requirements too...

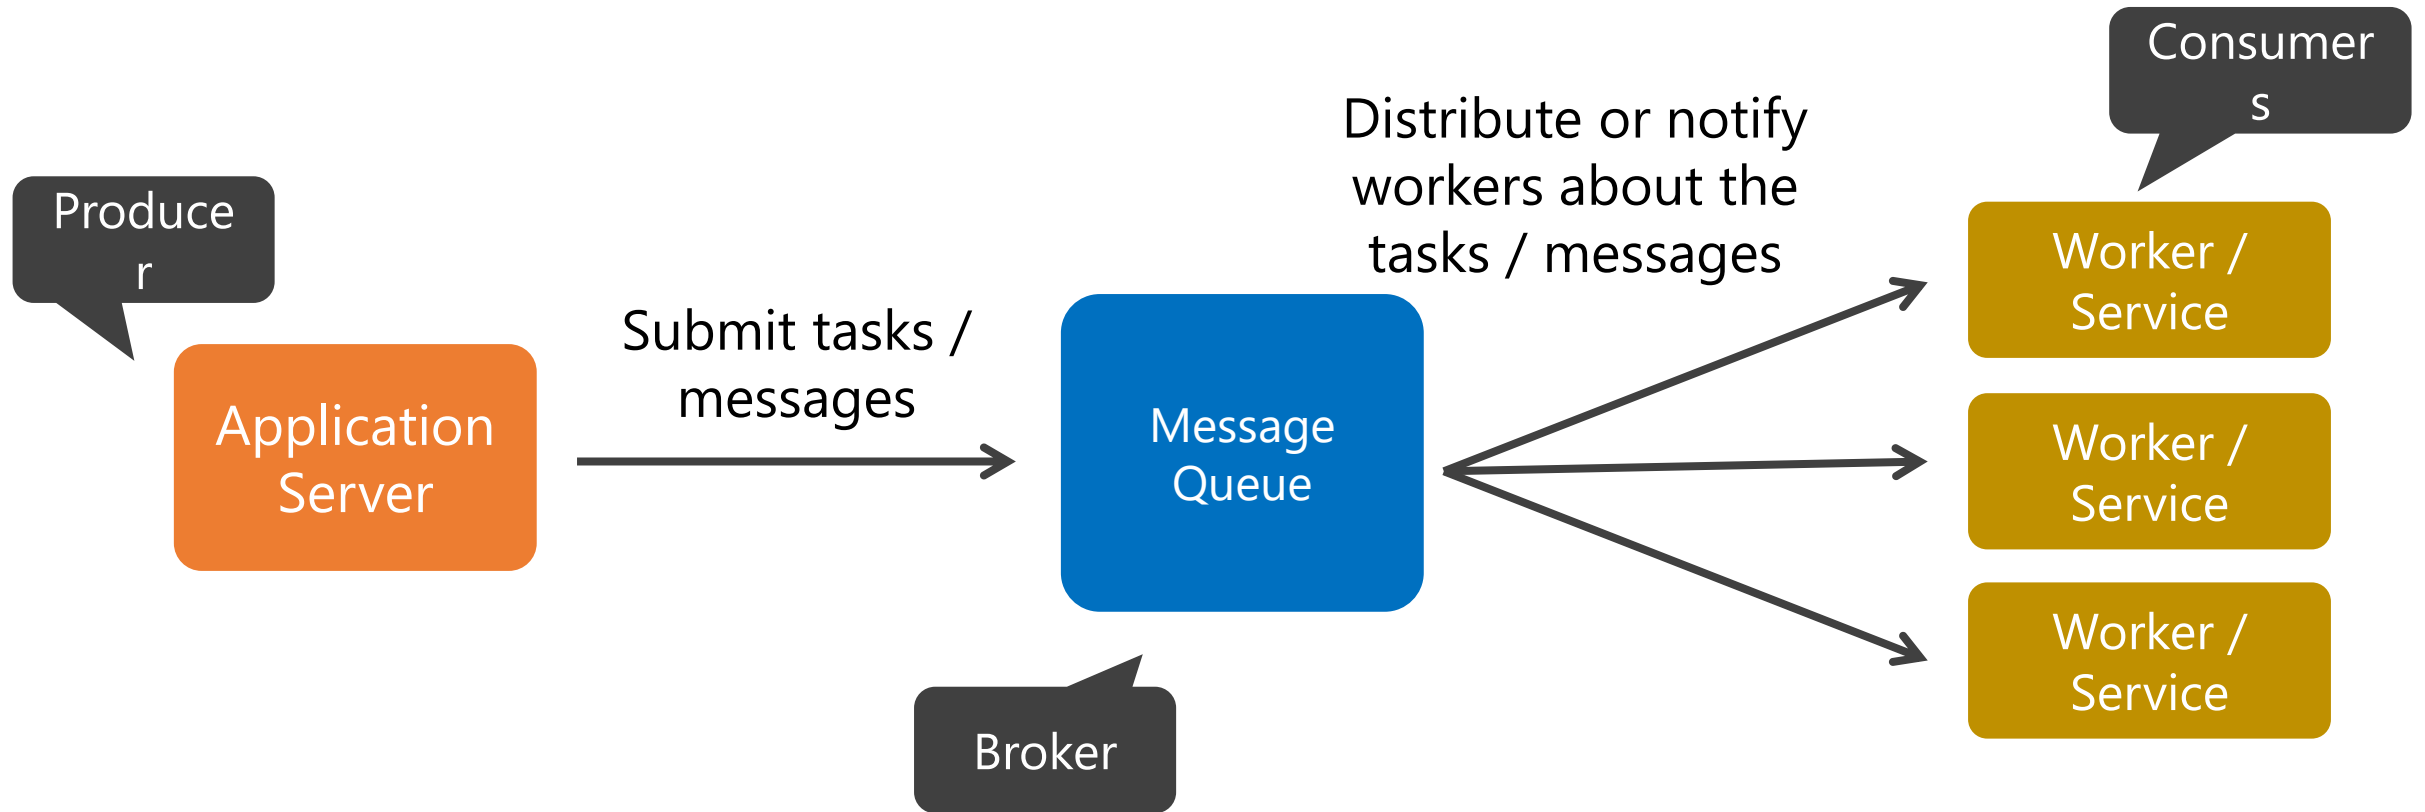
# Asynchronous Tasks

It can be non-trivial to implement asynchronous tasks:

- Tasks may have to be executed **sequentially**
- There can be a lot of **concurrent** tasks
- Tasks may have different **priorities**
- Different tasks may require different **amount of resources**

We need some "**middle man**" to help us manage the tasks/messages between the server application and the workers

# Message Queues



# Message Queues

By having a “broker” between the application server and the services, we actually have a more **robust** system

- Free the HTTP request-response cycle from heavy tasks
- Clients are shielded from failures of background tasks
- If there is a failure, the broker can make sure that the task is submitted again for retry

# Asynchronous Message Queues

A popular way of implementing asynchronous message queue in Python applications involves using the combination of the following:

- Celery – as the task queue
- RabbitMQ – as the message broker
- Redis – as the backend data store

# RabbitMQ

# RabbitMQ

One of the commonly used open source message broker software. Some features include:

- Delivery acknowledgement
- Clustering and queues mirroring
- Client libraries available in many programming languages
- Management UI

References:

<https://www.rabbitmq.com/>

<http://blogs.vmware.com/vfabric/2013/03/scaling-real-time-comments-huffpost-live-with-rabbitmq.html>

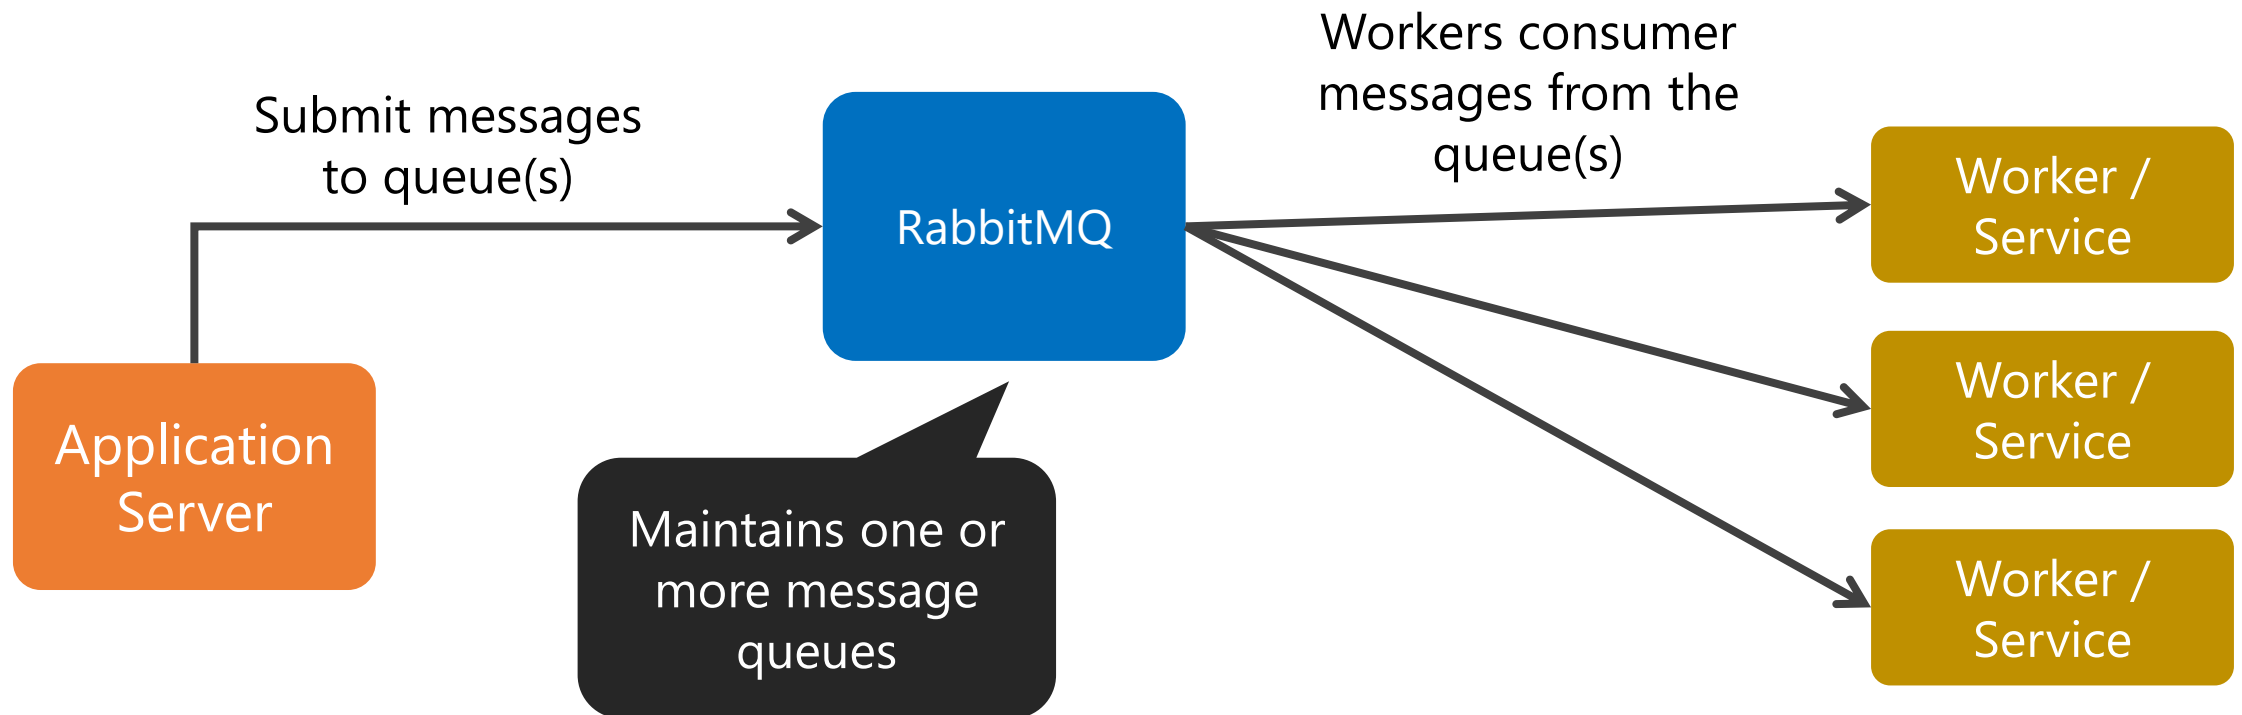
<http://blogs.vmware.com/vfabric/2013/03/how-indeed-com-handles-35-million-job-postings-per-day-using-rabbitmq.html>





# RabbitMQ

Architecture when using RabbitMQ as the message broker (Note: the components below do not necessarily reside on the *same machine*)



# RabbitMQ

For installation, refers to <https://www.rabbitmq.com/download.html>

RabbitMQ uses the AMQP protocol for message passing

- AMQP = *Advanced Message Queuing Protocol*
- An application layer protocol for sending and receiving messages
- Defines how messages are routed and stored in the broker
- Defines how communications are done between clients and server (broker)

# RabbitMQ – Sending Messages

In Python, we can install a module that help us use AMQP to talk to RabbitMQ (e.g. pika)

The block on the right shows an example of submitting a message to a queue in RabbitMQ

```
import pika

connection = pika.BlockingConnection(
    pika.ConnectionParameters('localhost'))
channel = connection.channel()

channel.queue_declare(queue='queue001')

channel.basic_publish(
    exchange='',
    routing_key='queue001',
    body='This is a message!')

connection.close()
```

Put this in your  
server  
application

Ref: <https://www.rabbitmq.com/tutorials/tutorial-one-python.html>

# RabbitMQ – Consuming Messages

Consuming a message requires more complex codes, as it involves creating a *callback function* to handle incoming message

```
import pika

connection =
    pika.BlockingConnection(pika.ConnectionParameters(host='localhost'))
channel = connection.channel()
channel.queue_declare(queue='queue001')

def callback(ch, method, properties, body):
    print "Received %r" % (body,)

channel.basic_consume(callback, queue='queue001', no_ack=True)
channel.start_consuming()
```

Execute this script  
on the worker  
machine

Ref: <https://www.rabbitmq.com/tutorials/tutorial-one-python.html>

# RabbitMQ

The above example refers to the most basic usage of RabbitMQ

- A single queue with a single consumer

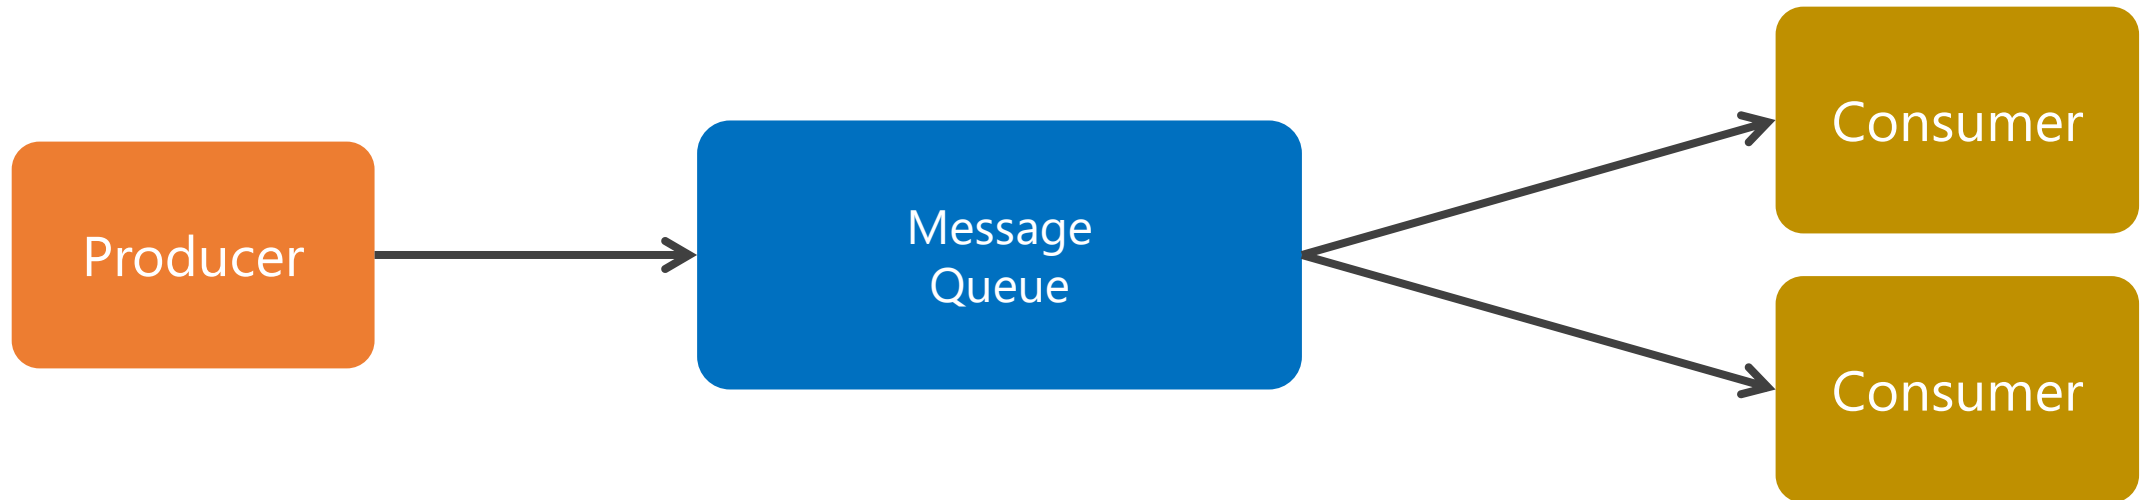


*What if we want to have more consumers to share the load, introduce different types of messages, ...?*

# RabbitMQ

Distributing messages among multiple workers


- Simply create more workers and make them consume *from the same queue*
- RabbitMQ will distribute messages to them in a *round-robin* fashion



# RabbitMQ

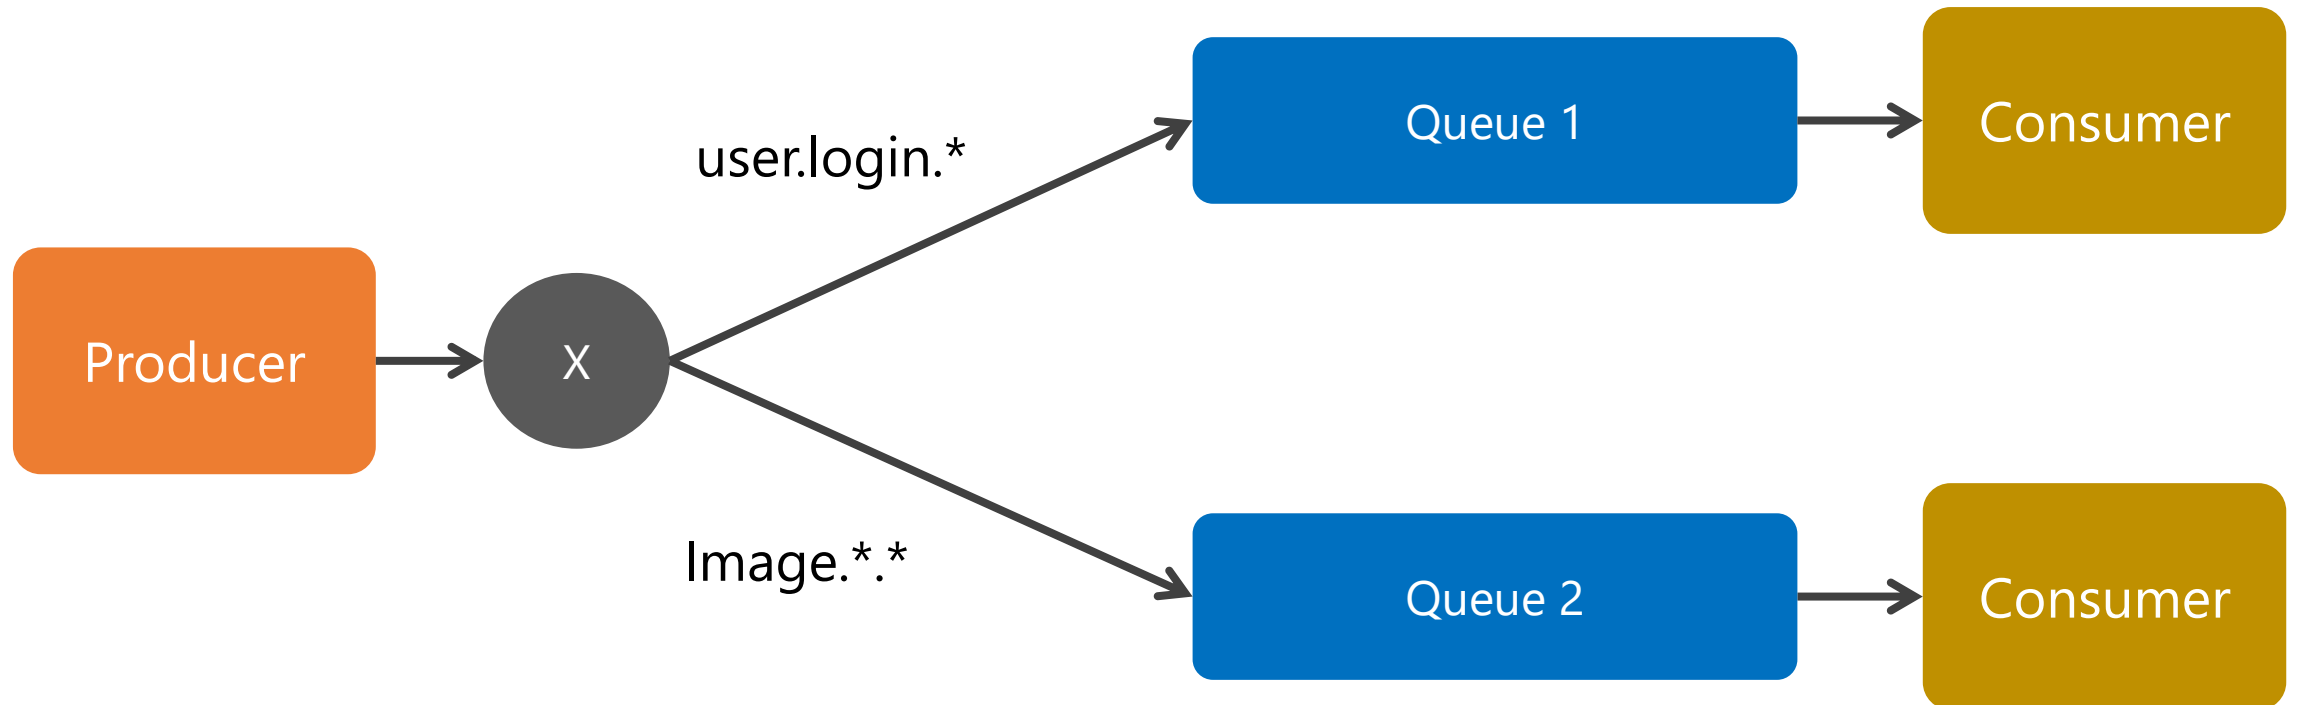
Broadcast message to all consumers at once (**publish/subscribe model**)

- Setup an **exchange** and let the exchange submit messages to multiple queues for multiple consumers



# RabbitMQ

Messages can also be routed by the exchange based on their “**topics**”, e.g.:




Ref: <https://www.rabbitmq.com/tutorials/tutorial-five-python.html>



# RabbitMQ

RabbitMQ is a general purpose message broker, you can implement the producers and consumers using different technology and languages (as long as you use the AMQP protocol)



## Events vs. Commands

The content of the message can be customised to your needs

- You can issue "**commands**". E.g. increment user's number of likes
- You can also issue "**events**". E.g. user ID=3 likes photo ID=5
- In general, it is better to send messages of "events"  
*(make your application event-driven!)*
  - › Events can be consumed by workers responsible for doing different tasks
  - › Allow better isolation between systems
  - › Do not need to hard-code actions in your application

# Celery

# Celery

- A **distributed task queue written in Python** for Python apps
- It has to be supported by a message broker (e.g. **RabbitMQ**)
- Allow implementation of asynchronous tasks to be more integrated into your Python application

*Confused?*

Ref: <http://www.celeryproject.org/>

# Celery

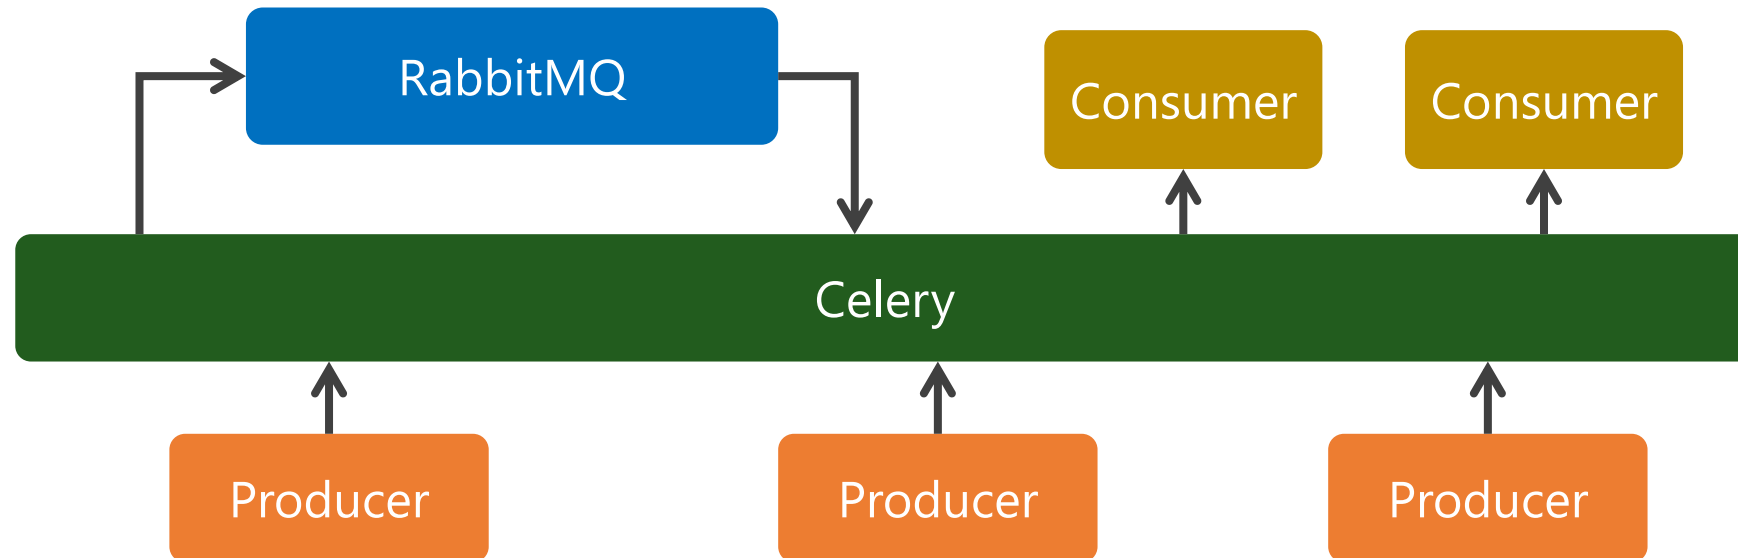
Before you use Celery:



# Celery

After you use Celery

- Allows you to send messages / invoke asynchronous tasks by simply making a function call in Python
- Helps you manage your workers (e.g. restart them in case of a failure or if an exception is raised)



# Celery – Example

Let's start from defining the asynchronous task you want to perform

Give the module a name (e.g. 'tasks')

Specify the URI of the RabbitMQ server

```
from celery import Celery

app = Celery('tasks', broker='amqp://guest@localhost//')

@app.task
def countWordsInWebPage(url):
    content = downloadURLContent(url)
    words = countWords(content)
    hash = updateDatabase(words)
    return hash
```

Define a task

## Celery – Example

After defining the tasks, run the Celery worker server with:

```
$ celery -A tasks worker
```

Start Celery,  
which will host  
the worker

And then you can perform the defined tasks asynchronously in your application:

```
from tasks import countWordsInWebPage  
url = "https://en.wikipedia.org/wiki/RMS_Titanic"  
countWordsInWebPage.delay(url)  
...
```

Import the  
function from  
your task module

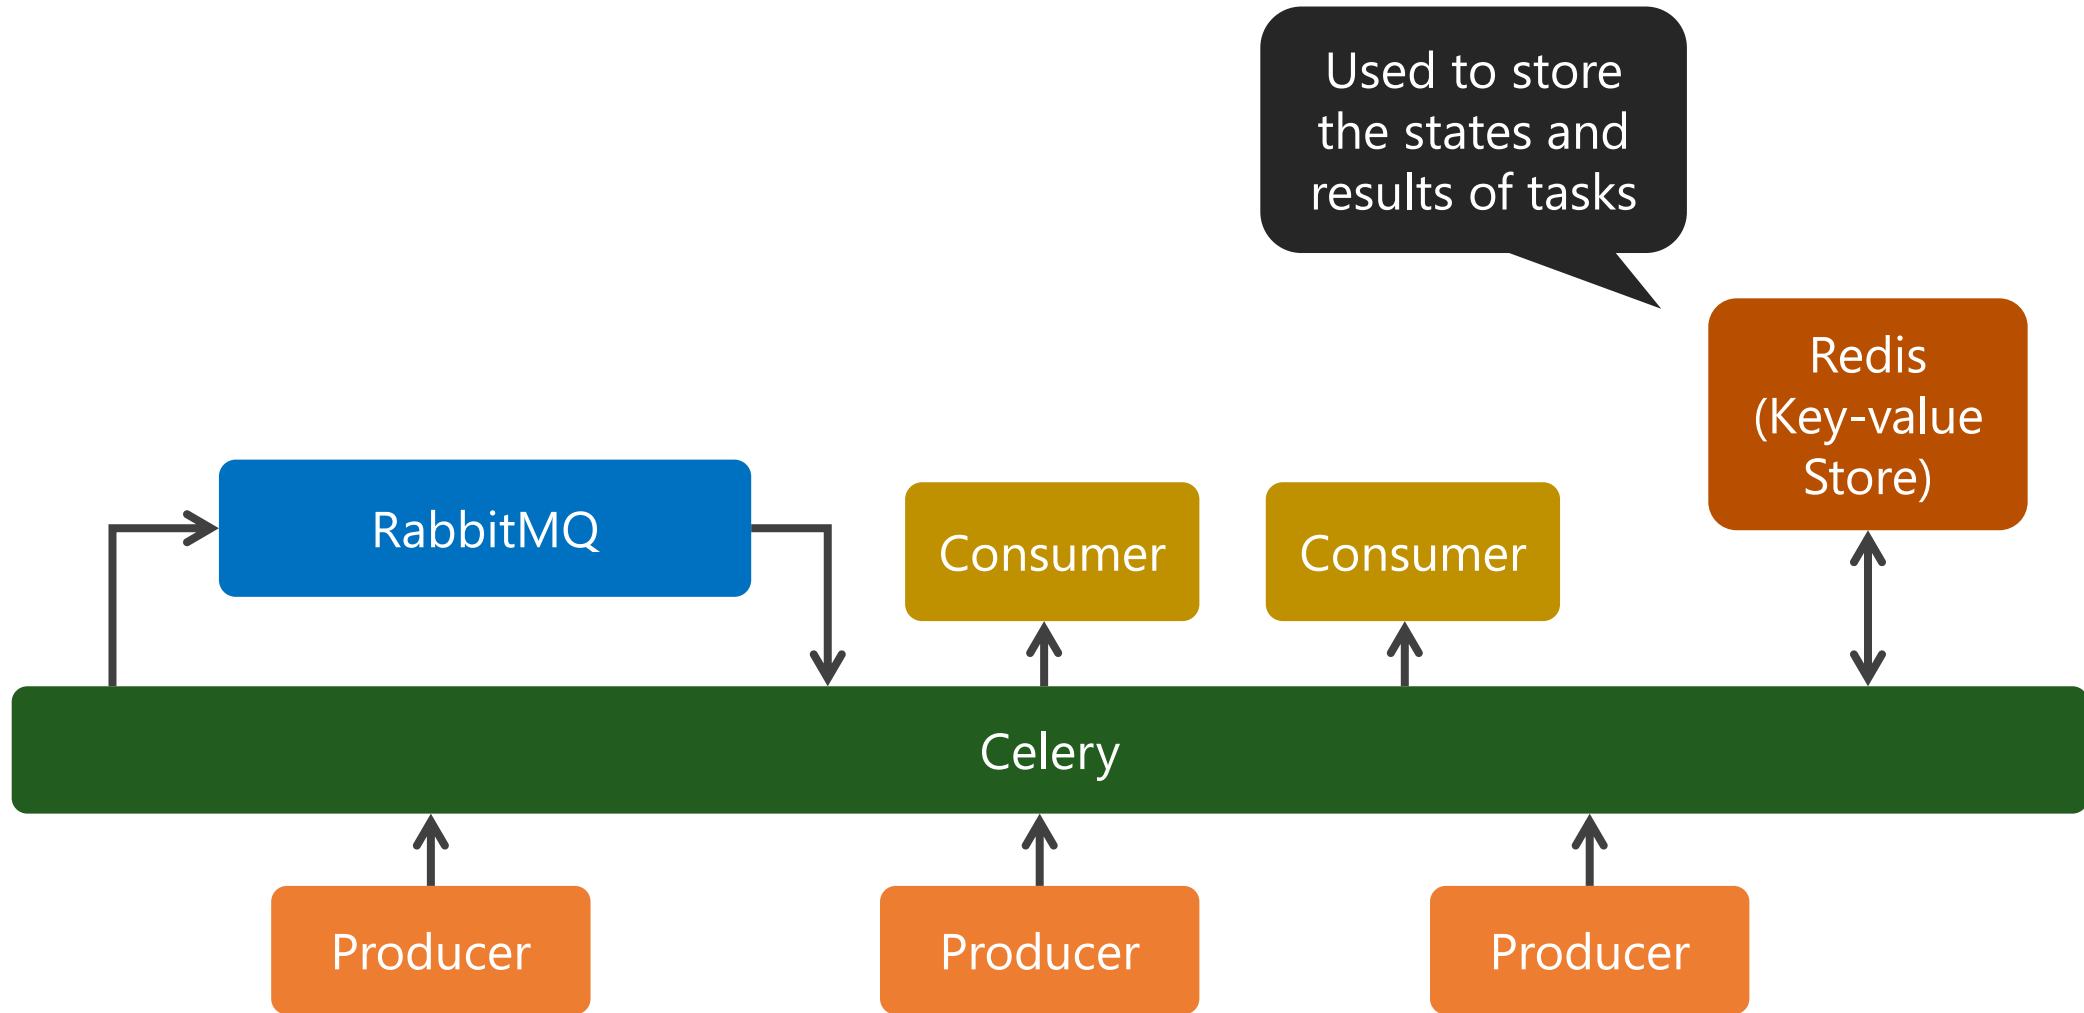
Invoke the 'delay'  
method to perform the  
task asynchronously



# Celery

- In many cases, you simply want to submit a task and are not concerned about the result (e.g. user likes and article, user comments on a photo, etc.)
- In other cases, you may want to *keep track of the status of the task* (e.g. let the user know about the progress of uploading a file)
- For the latter case, Celery needs a “**backend**” **storage** to temporarily stores the states of the asynchronous tasks
- Usually **Redis** is used as the backend

# Celery



# Celery

## Setting the backend of Celery:

```
from celery import Celery

app = Celery('tasks',
             backend='redis://localhost', broker='amqp://guest@localhost//')

@app.task
def countWordsInWebPage(url):
    content = downloadURLContent(url)
    words = countWords(content)
    hash = updateDatabase(words)
    return hash
```

# Celery

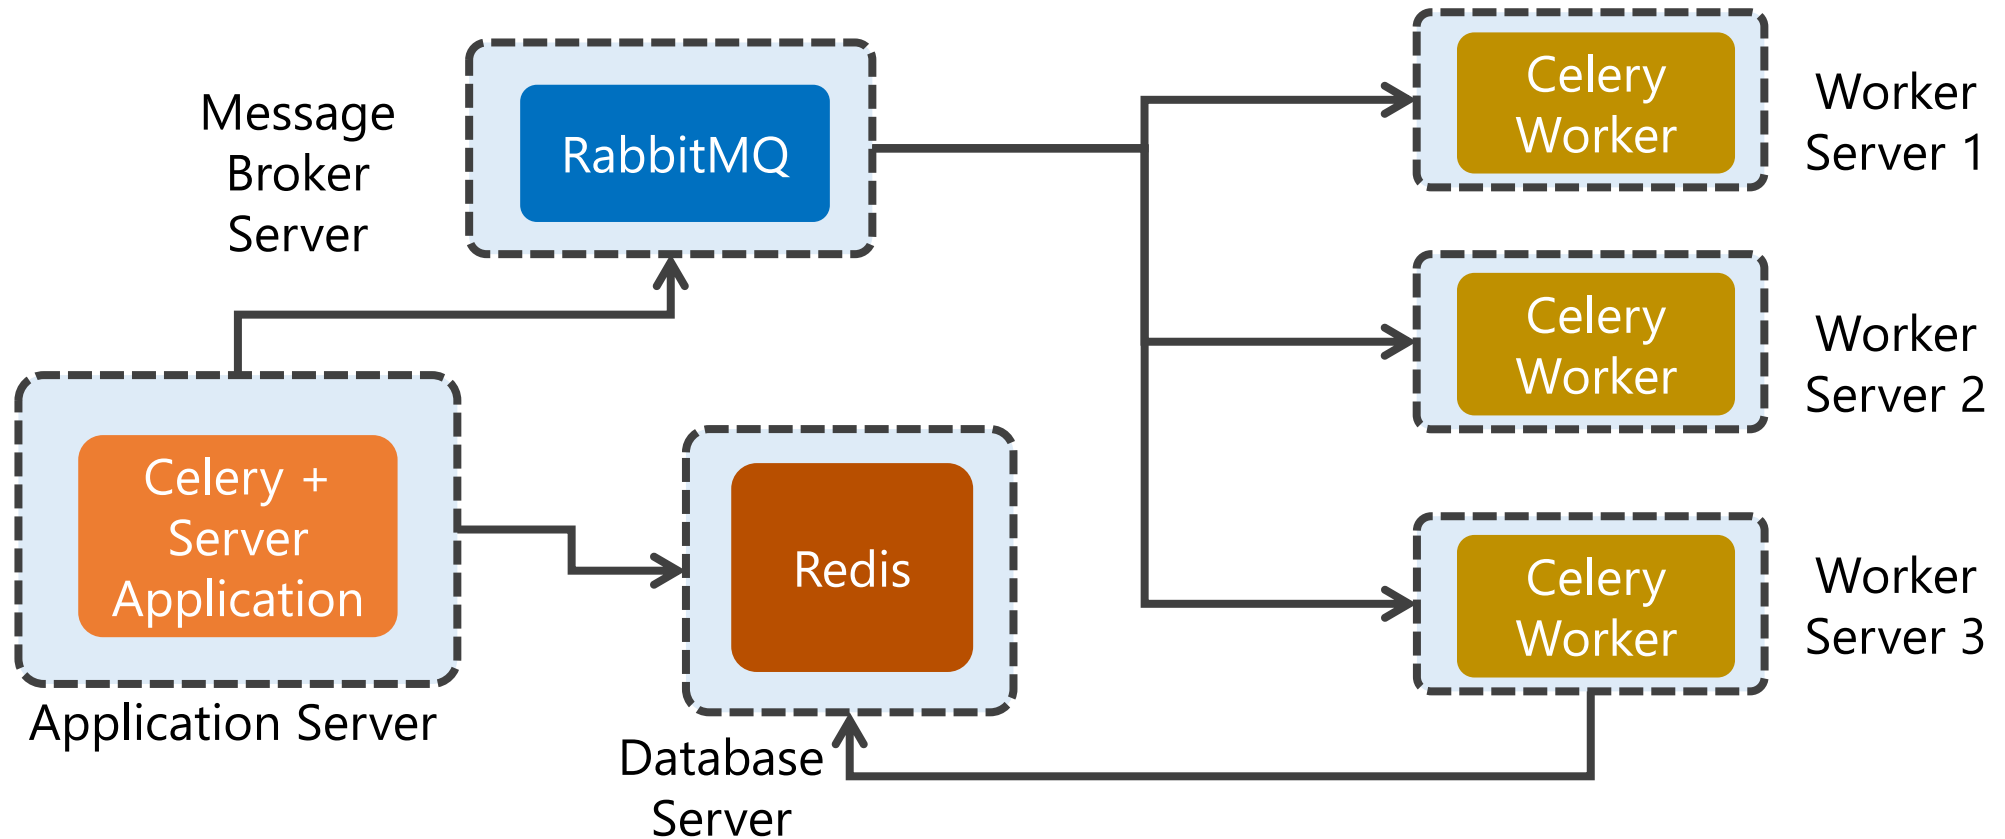
Once you have a backend, you can check the status of a task submitted:

```
...  
result = countWordsInWebPage.delay(url)  
task_id = result.task_id  
...
```

```
from celery.result import AsyncResult  
result = AsyncResult("my-task-id")  
result.ready() # True if the task is finished  
result.result  # The return value of the task  
result.state   # The current state of the task
```

# Celery

The components can all be set on different machines for scalability  
(Workers should be pointed to the URI of the same message broker)



# Celery + Flask

You can easily integrate Celery with your Flask application  
Firstly, define a function that creates a new Celery object:

```
from celery import Celery

def make_celery(app):
    celery = Celery(app.import_name,
broker=app.config['CELERY_BROKER_URL'])
    celery.conf.update(app.config)
    TaskBase = celery.Task

    class ContextTask(TaskBase):
        abstract = True
        def __call__(self, *args, **kwargs):
            with app.app_context():
                return TaskBase.__call__(self, *args, **kwargs)
    celery.Task = ContextTask
    return celery
```

# Celery + Flask

Then, create a Flask app and use it to initialise the Celery object and define some tasks using the **@celery.task()** decorator

```
from flask import Flask

app = Flask(__name__)
app.config.update(
    CELERY_BROKER_URL='amqp://guest@localhost',
    CELERY_RESULT_BACKEND='redis://localhost:6379'
)
celery = make_celery(app)

@celery.task()
def add(a, b):
    return a + b
```

# Celery + Flask

Then, you can create asynchronous tasks by invoking the functions in your application. For example:

```
@app.route('/')  
def do_add():  
    add.delay(2,3)  
    return "Done"
```

Ref:

<http://flask.pocoo.org/docs/0.10/patterns/celery/>

<https://github.com/miguelgrinberg/flask-celery-example>



# Celery + Flask

Before deploying your Flask app, you will also have to start the worker using Celery, for example:

```
$ celery -A app.celery worker
```

If you need more debug information, use the `--loglevel` parameter:

```
$ celery -A app.celery worker --loglevel=DEBUG
```

Ref:

<http://flask.pocoo.org/docs/0.10/patterns/celery/>

<https://github.com/miguelgrinberg/flask-celery-example>

# Asynchronous Message Queues + NoSQL Database

## Case Study

Consider a *social application*, in which users can be *friends* of each other.

What are some of the common operations?

- Get a list of friends of a given user
- Get a list of common friends of two given users
- Recommend a user add other users as friends if they are friends of his/her friends
- ...

# Case Study

A straight-forward way of storing the friendship network is to use a table with the following schema in a relational database:

**TABLE friendship**

Column	TYPE
user_id	INTEGER
friend_id	INTEGER
is_friend_since	DATE

**Data**

user_id	friend_id	is_friend_since
1	2	2015-05-07
2	1	2015-05-07
3	2	2015-09-20
2	3	2015-09-20
...	...	...

## Case Study

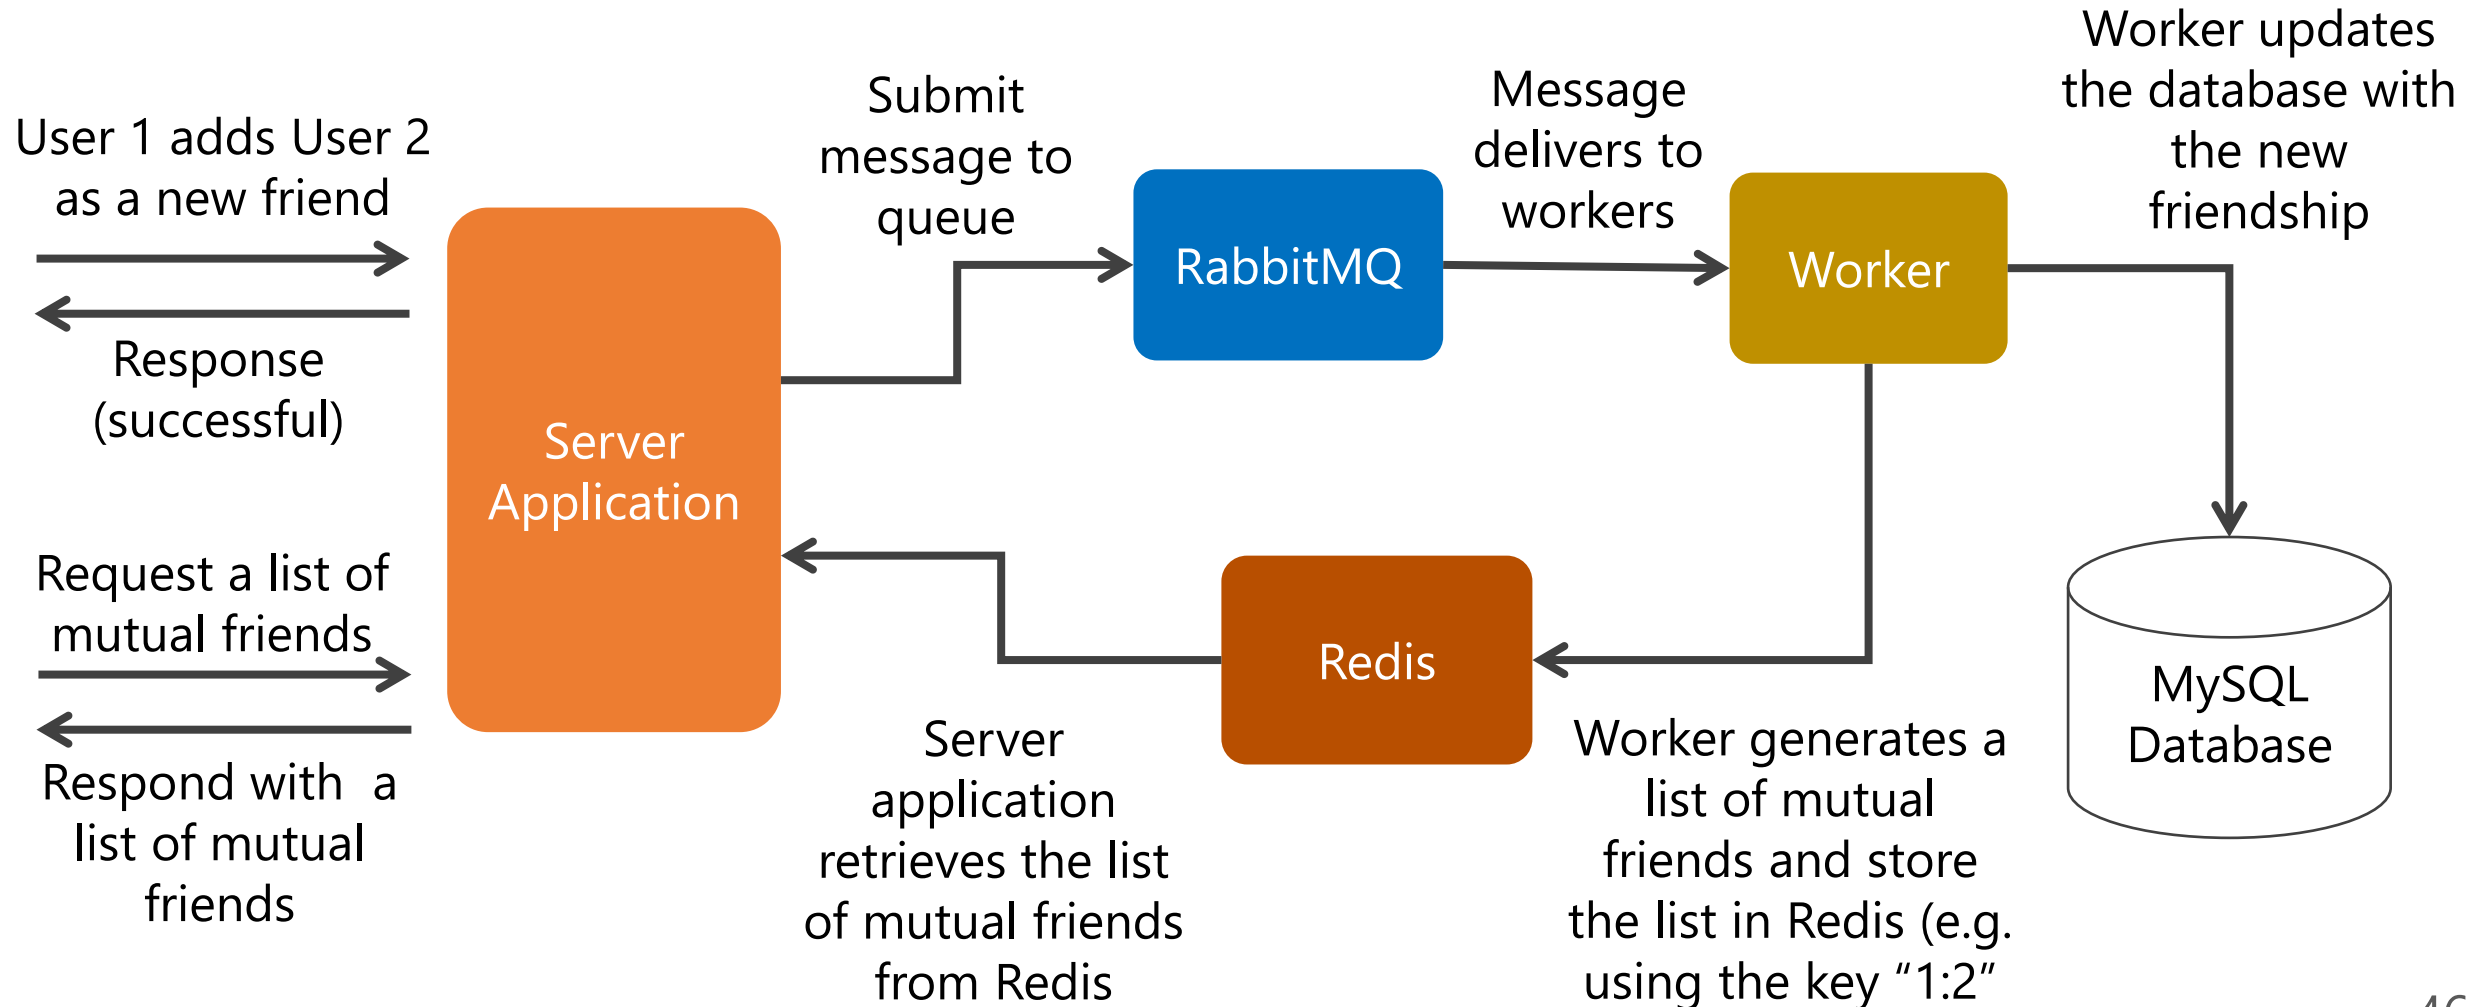
How do you get **a list of mutual friends** of two given users (say user 2 and user 3)?

- Retrieve the lists of friends of these two users, and then find the overlap of these two lists in your application code
- Perform a complex SQL query by joining the friendship table with itself and let the database return the list of mutual friends

It would not be efficient to compute this list on-the-fly. It is also likely that this list is not going to be changed very frequently.

# Case Study

How can we improve the efficient of this function using asynchronous tasks?



End of Lecture 9